

1.W ODI Data Ethics Canvas Exercise

Resources

- [ODI Data Ethics Canvas](#)
- [ODI Data Ethics Canvas User Guide](#)
- [Ethics First AI.png](#) (in your GitHub repo!)

Scenario

Your team has just run the QuickLoan pipeline. Before it goes anywhere near production, you're the governance team doing a risk check.

You'll each take on a role. Pick one or assign them:

- **AI Engineer** – thinking about technical debt, model drift, and security
- **Compliance Lead** – thinking about GDPR, FCA, and audit trails
- **Product Manager** – thinking about customer trust and reputation

Your job is to **fill out the canvas**, using the pipeline (and CRISP-ML(Q)) as your lens. Think about each box in the canvas across the whole lifecycle, from data collection to deployment.

Try to fill out the entire canvas. Be thoughtful, but decisive. You don't need to be perfect.

As you go, ask yourselves:

- Who could be harmed by a bad decision?
- What biases might be baked into our data?
- What power dynamics are at play here?
- What trade-offs are we making: speed vs fairness, profit vs trust?

You can record some of your thoughts in your lessons learnt log. Section 3 is most relevant here.

Output

Each team will share **one key risk or trade-off** that you think absolutely must be addressed before this model goes live. That means one person in your team will be sharing, so think about who's most comfortable doing that, and maybe decide that while you're assigning your roles.